

Naman Padiyar

naman.padiyar09@gmail.com | (+91)7453030070 | [LinkedIn](#) | [LeetCode](#) | [GitHub](#)

Education

Delhi Technological University

Bachelor of Technology in Computer Science

Nov 2022 – May 2026

9.00 GPA

Adarsh J Dharmic Shiksha Sadan

CBSE Class XII (Science)

Apr 2021 – Jul 2022

92.8%

Parvati Radhakishan Fomra School

CBSE Class X

Apr 2019 – Jul 2020

94.33%

Experience

Software Development Engineer Intern

Innovaccer

Jan 2026 – Present

Noida, India

- Built an end-to-end **AI-native Deal Desk platform** spanning **5 healthcare product brands**, translating product requirements into production Discovery, Pricing, and Margin workflows for Deal Desk teams.
- Engineered an **LLM-driven Discovery agent** that auto-answers an **83-question** deal discovery framework with cited evidence from emails, call notes, and legal docs via **MCP** tool-use – cutting a multi-hour manual review into a **~60-second** automated run.
- Collaborated cross-functionally with engineers and product teams to integrate AI agents into existing enterprise systems, reducing manual intervention in compliance workflows and improving deal-cycle operational efficiency.
- Used **AI coding assistants** (Claude, GitHub Copilot) to accelerate development, testing, and iteration on document processing pipelines, while maintaining code review and engineering standards.

Projects

Enterprise KnowledgeHub AI – Hybrid RAG Platform | [GitHub](#)

- Architected a full-stack **Hybrid RAG** platform (React, TypeScript, FastAPI, MongoDB) using **FAISS** + **BM25** retrieval (**384-dim embeddings**), **RRF** fusion, and cross-encoder reranking for citation-backed Q&A.
- Orchestrated agentic reasoning via an **8-node LangGraph** pipeline (confidence-gated retrieval, live **Tavily** web-search fallback) with full-trace observability via **LangSmith**.
- Implemented async document ingestion (OCR, image captioning) on a **Redis**-backed queue; **Redis** caching cut repeat-search latency **50x** (**4.1s** → **0.08s**); found and fixed **2 production bugs** via live end-to-end verification.

GAN-Based Steganography System | [GitHub](#)

- Designed a multi-modal GAN-based system to imperceptibly hide secret data within Image, Audio, and Video files; deployed as a containerized **FastAPI** service with **AES-256-GCM** encryption.
- Evaluated model performance across standard quality metrics; integrated a deep learning-based text compression module reducing payload by **40–60%**, benchmarked against classical methods.

AI-Powered Job Portal | [GitHub](#)

- Developed a responsive web interface for a **Gemini AI-powered** resume-JD semantic matching engine, presenting NLP-based candidate scoring and ranking across 50+ job applications through a clean, filterable UI.
- Designed the end-to-end user flow from resume upload to real-time ranked results, deploying via **Vercel** with **Sentry** for real-time error tracking and UI reliability monitoring in production.

Technical Skills

Frontend: JavaScript, TypeScript, React, HTML/CSS, REST API Integration, Responsive UI, Client-side Validation

Languages: C++, C, Python, SQL

Fundamentals: OOP, DSA, Data Structures, Operating Systems, Computer Networks, DBMS

Tools/Other: Git, AWS S3, FastAPI, Docker, MongoDB, Redis, LangChain, LangGraph, Claude, Salesforce, PyTorch, Scikit-learn, Pandas, NumPy

Achievements

- Authored a **research paper** on adversarial GAN-based steganography with spatio-temporal modeling for multi-modal secret data embedding across image, audio, and video.
- Secured top ranks in **JEE Mains 2022 (AIR 6172, 99.33 percentile)** and **JEE Advanced 2022 (AIR 6417)** among 1M+ candidates.
- Solved **800+** coding challenges on LeetCode and CodeChef, earning **Knight** rank (**1800+ rating**) for proficiency in complex data structures.